

# CS307

## Database Principles

Stéphane Faroult  
faroult@sustc.edu.cn

Liu Zijian      liuzijian47@163.com

MidTerm Exam: 10<sup>th</sup> November

Not what I said at  
the beginning of  
the lecture

Next Week: Review

As we have seen last time, when you use technical, integer identifiers as primary keys (usually to replace a combination of several unique columns that are the "real life" key) you cannot simply find the highest current value and add one, because on a busy database adds are very high that two sessions will read the same value – and try to insert the same value. It will work for one, and the second one will get a constraint violation error. The solution is to let the system manage the generation of new identifiers, and there are two ways of doing it.

## SEQUENCE

**create sequence movie\_seq**

You can use special database objects called sequences, which are simply number generators. By default they start with 1 and increase by 1 (they can reach values that are very, very big)



**ORACLE**

```
insert into movies(movieid, ...)
values(movie_seq.nextval, ...)
insert into credits(movieid, ...)
values(movie_seq.currval, ...)
```

**IBM DB2**

```
insert into movies(movieid, ...)
values(next value for movie_seq, ...)
insert into credits(movieid, ...)
values(previous value for movie_seq, ...)
```

**PostgreSQL**

```
insert into movies(movieid, ...)
values(nextval('movie_seq'), ...)
insert into credits(movieid, ...)
values(currval('movie_seq'), ...)
```

Syntax varies, but you can obtain a new (guaranteed to be unique) number, and retrieve the last number you obtained for this sequence and this session.

DB2 only

## AUTO-NUMBERED COLUMN

The other approach is to give a special definition to a column saying that it will automatically get increasing values (you are usually limited to one such column per table)



## AUTO-NUMBERED COLUMN

**create table movies**

**Microsoft SQL Server** (movieid int not null identity primary key,

**PostgreSQL** serial primary key,

**MySQL** integer primary key,

**SQLite**

MySQL can also use an AUTOINCREMENT attribute.

As usual, syntax differs. PostgreSQL actually creates a sequence behind the scene, which it "attaches" to the table so that dropping the table drops the sequence.

## AUTO-NUMBERED COLUMN

Oracle (since version 12, it wasn't possible before) can do it PostgreSQL style, but more explicitly.

define

movieseq.nextval

as default value for movieid

Note that if you drop the table, the sequence will still exist.

**ORACLE** >= 12c

If you insert a film with an auto-numbered column, you just omit the movieid from the INSERT statement, it will get automatically populated.

```
insert into movies(title, ...)  
values('Some Movie Title', ...)
```

To retrieve the last value generated in your session, you use a special variable such as @@identity with SQL Server, or functions with other products (eg. lastval() with PostgreSQL or last\_insert\_id() with MySQL). This last value is often "across all tables", although sometimes you can retrieve the last value for a specific table.

First Name

Last Name

Email

Gender --- ▼

Born ---- ▼

Usually in interactive programs, interfaces collect data for one row and issue the corresponding INSERT statement that adds the row to the right table.

**insert**

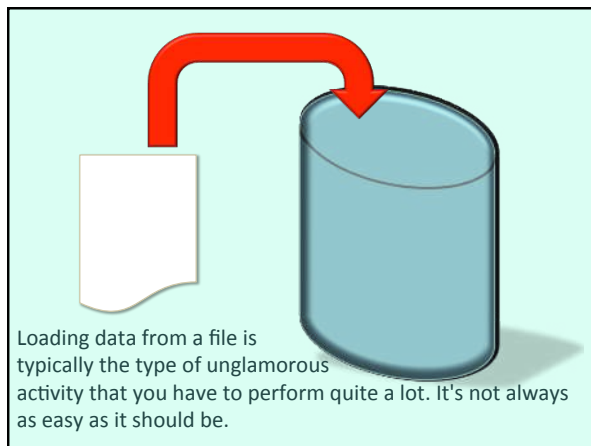
However, it's extremely common that you want to upload data in bulk, either to populate a database initially or because data is exchanged between different systems by using files.



Flickr: Elbfoto


```
insert into table_name  
(column1, column2, ..., columnn)  
select col1, col2, ..., coln  
from ...
```


Another way to massively insert data in a table is by inserting the result of a query. Of course, it assumes that the data is already in the database! What usually happens is that data is loaded "as is" in work tables (staging areas) that are badly normalized (they are hardly more than the table image of a file) then dispatched through INSERT ... SELECT ... statements to the well designed tables.



**Line 1**  
**Line 2**

First of all, when you have two lines in a text file, they are separated by one character if the file comes from a Linux system, two characters if it comes from a Windows system. When you transfer data from one to the other you need to be careful.

**Line 1** **\n** **Line 2**  


**Line 1** **\r\n** **Line 2**  


**CSV**

*title* *year* *minutes* *Black&white* *Color*

"Citizen Kane",	1941,	119,	B
"The Godfather",	1972,	175,	C
"Taxi Driver",	1976,	113,	C
"Casablanca",	1942,	102,	B
"Raging Bull",	1980,	129,	C
"Singin' in the Rain",	1952,	103,	C
"North By Northwest",	1959,	136,	C
"Gone with the Wind",	1939,	226,	C

A very popular format is the Comma Separated Values format, in which some fields may be enclosed by double quotes.

us\_movie\_info.csv

**Tab-separated**

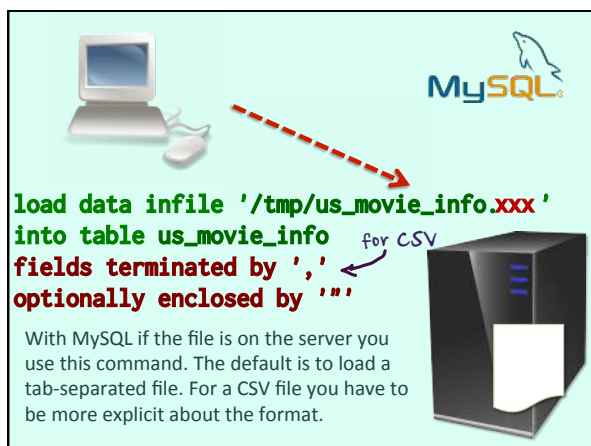
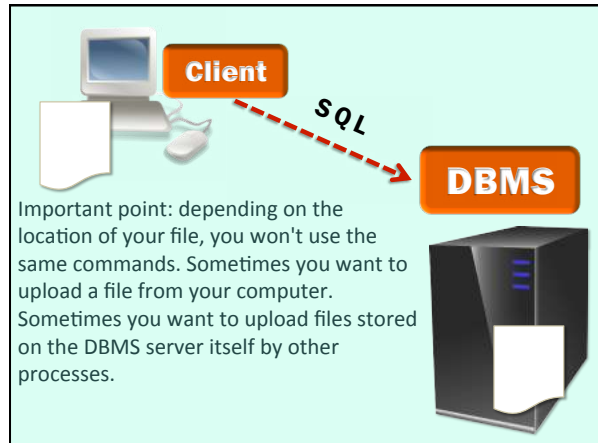
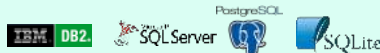
Another popular format is to have tab-separated fields.

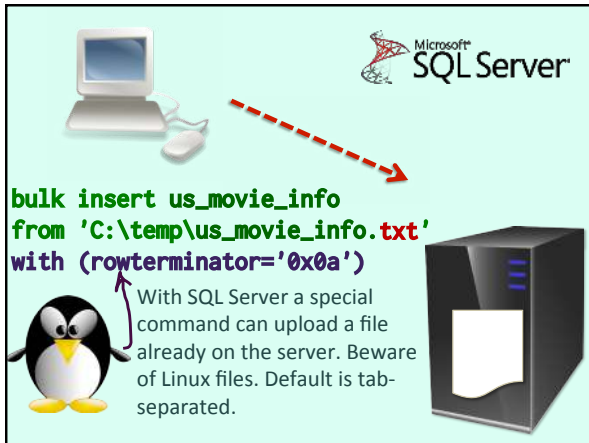
Citizen Kane	1941	119	B
The Godfather	1972	175	C
Taxi Driver	1976	113	C
Casablanca	1942	102	B
Raging Bull	1980	129	C
Singin' in the Rain	1952	103	C
North By Northwest	1959	136	C
Gone with the Wind	1939	226	C

us\_movie\_info.txt

Let's create a "staging table" with one column to receive every field in the file. This syntax wouldn't work with MySQL (wants DECIMAL instead of NUMERIC), nor with Oracle (wants NUMBER).

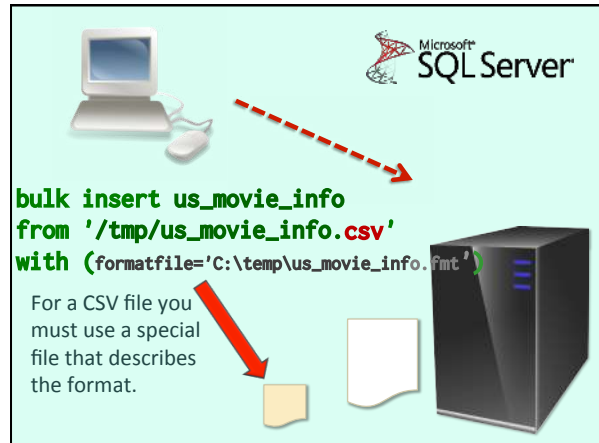
```
create table us_movie_info
(title          varchar(100) not null,
year_released  numeric(4) not null,
duration       int not null,
color          char(1) not null,
primary key(title, year_released))
```





**bulk insert us\_movie\_info  
from 'C:\temp\us\_movie\_info.txt'  
with (rowterminator='0x0a')**

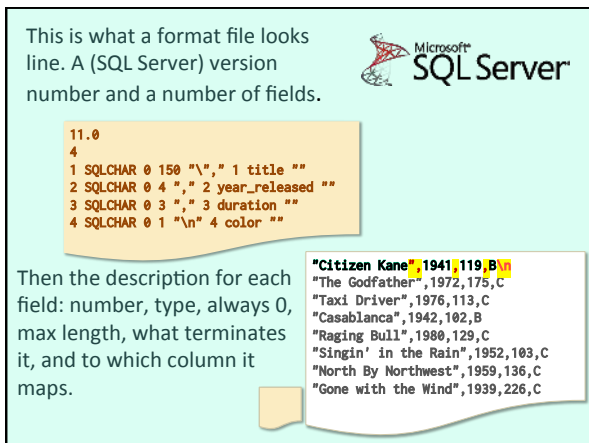
With SQL Server a special command can upload a file already on the server. Beware of Linux files. Default is tab-separated.



**bulk insert us\_movie\_info  
from '/tmp/us\_movie\_info.csv'  
with (formatfile='C:\temp\us\_movie\_info.fmt')**

For a CSV file you must use a special file that describes the format.

This is what a format file looks like. A (SQL Server) version number and a number of fields.

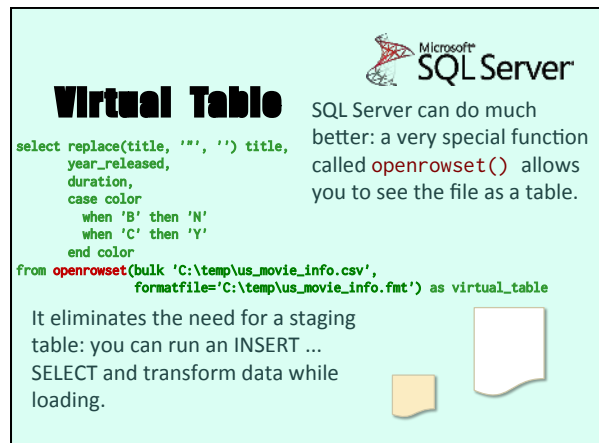


```
11.0
4
1 SQLCHAR 0 150 "\", " 1 title ""
2 SQLCHAR 0 4 ", " 2 year_released ""
3 SQLCHAR 0 3 ", " 3 duration ""
4 SQLCHAR 0 1 "\n" 4 color ""
```

Then the description for each field: number, type, always 0, max length, what terminates it, and to which column it maps.

```
"Citizen Kane",1941,119,B
"The Godfather",1972,175,C
"Taxi Driver",1976,113,C
"Casablanca",1942,102,B
"Raging Bull",1980,129,C
"Singin' in the Rain",1952,103,C
"North By Northwest",1959,136,C
"Gone with the Wind",1939,226,C
```

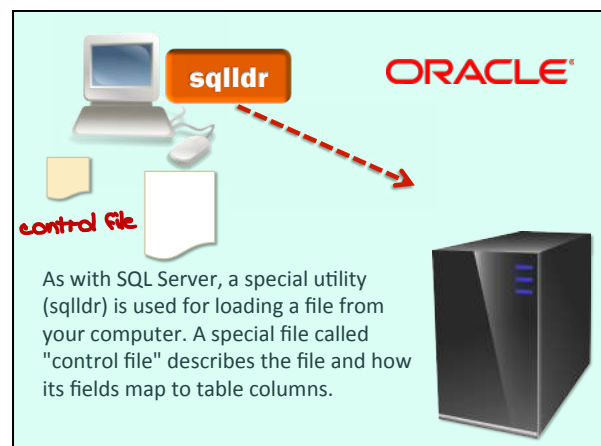
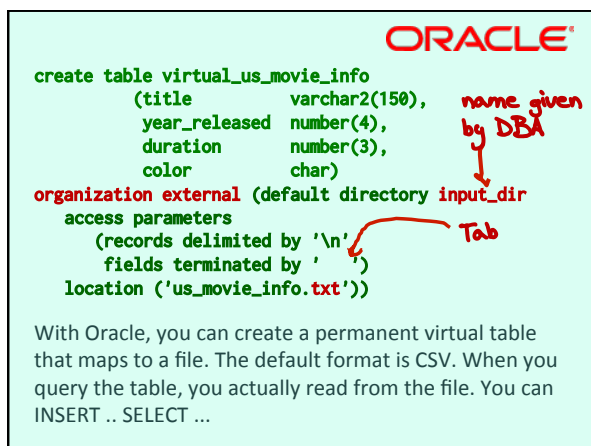
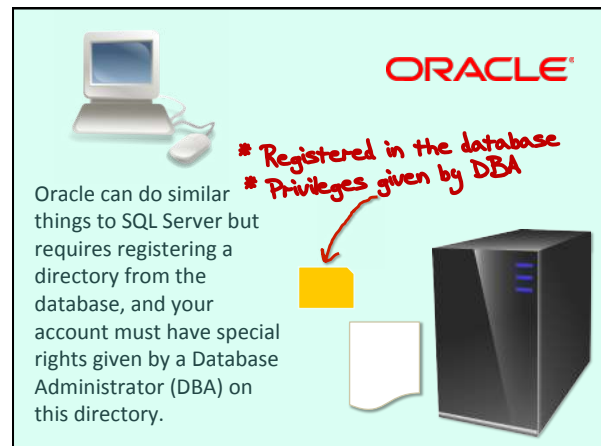
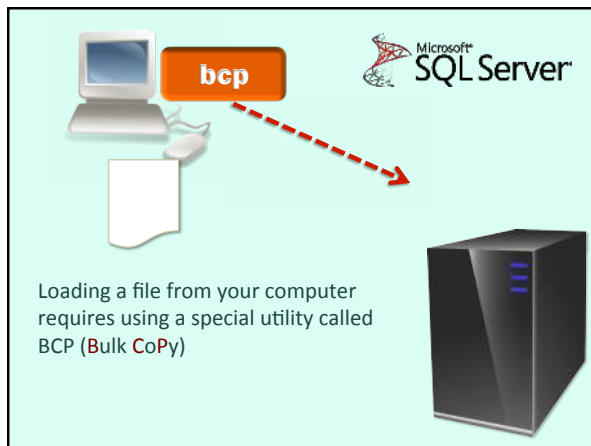
**Virtual Table**

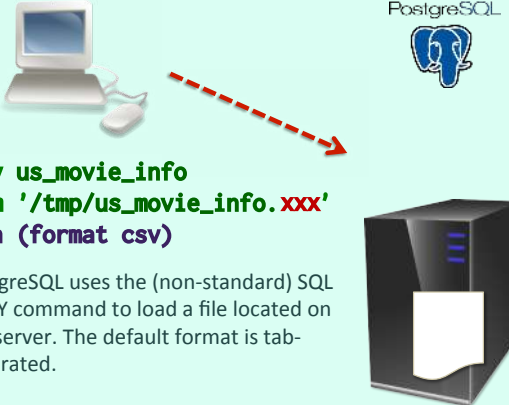


SQL Server can do much better: a very special function called **openrowset()** allows you to see the file as a table.

```
select replace(title, ',', '') title,
       year_released,
       duration,
       case color
         when 'B' then 'N'
         when 'C' then 'Y'
       end color
from openrowset(bulk 'C:\temp\us_movie_info.csv',
               formatfile='C:\temp\us_movie_info.fmt') as virtual_table
```

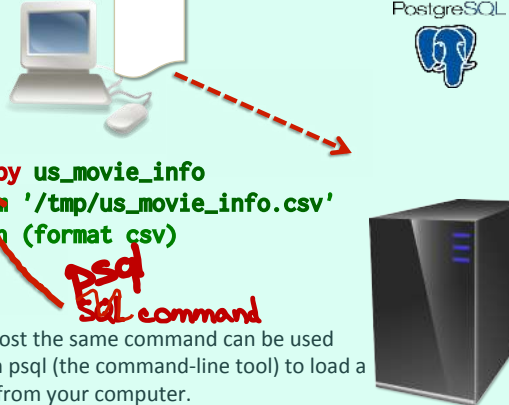
It eliminates the need for a staging table: you can run an INSERT ... SELECT and transform data while loading.





**copy us\_movie\_info  
from '/tmp/us\_movie\_info.xxx'  
with (format csv)**


PostgreSQL uses the (non-standard) SQL COPY command to load a file located on the server. The default format is tab-separated.



**\copy us\_movie\_info  
from '/tmp/us\_movie\_info.csv'  
with (format csv)**

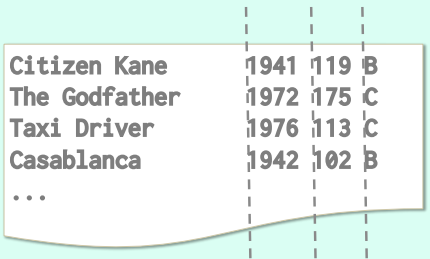
*psql  
SQL command*

Almost the same command can be used with psql (the command-line tool) to load a file from your computer.



**.separator ','  
.import 'us\_movie\_info.csv' us\_movie\_info**

You specify the separator, and load. It will trim double-quotes automatically.



Citizen Kane	1941	119	B
The Godfather	1972	175	C
Taxi Driver	1976	113	C
Casablanca	1942	102	B
...			


Although CSV and tab-separated files are most common, you can also encounter fixed-field files (fields always start at the same position), which all DBMS load easily.



```
<row><title>Citizen Kane</title><year>1941</year>...</row>
<row><title>The Godfather</title><year>1972</year>...</row>
<row><title>Taxi Driver</title><year>1976</year>...</row>
<row><title>Casablanca</title><year>1942</year>...</row>
...
```

XML is also a popular interchange format. The big products provide utilities for uploading XML files.

When everything else fails ...



Sometimes you encounter really weird text file formats. Using a scripting language to generate INSERT statements is usually the simplest solution.

# Things change ...



Source: Shutterstock

We have talked about inserting data, let's now see how we can update what is in the database.

Update is the command that changes column values. You can even set a non-mandatory column to NULL. The change is applied to all rows selected by the WHERE.

```
update table_name
set column_name = new_value,
    other_col = other_val,
    ...
where ...
```

```
update us_movie_info
set title = replace(title, '"', '')
```

Without a WHERE all rows are affected.

### Sorting issue with some names

We may want to modify some names in such a way as they sort as they should.

**D**

de Broca

These names should respectively appear under B, D and S.

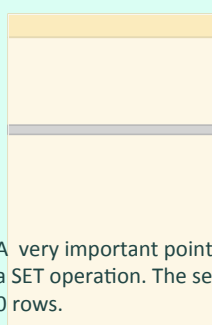
**V**

van Dijk

von Stroheim

```
update people
set surname = substr(surname, 4)
|| ' (von)'
where surname like 'von %'
```


This could be used to postfix all surnames starting by 'von' with '(von)' and turn for instance 'von Stroheim' into 'Stroheim (von)'



update

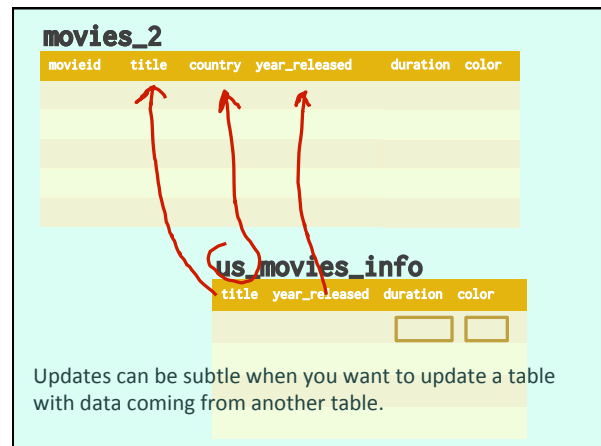
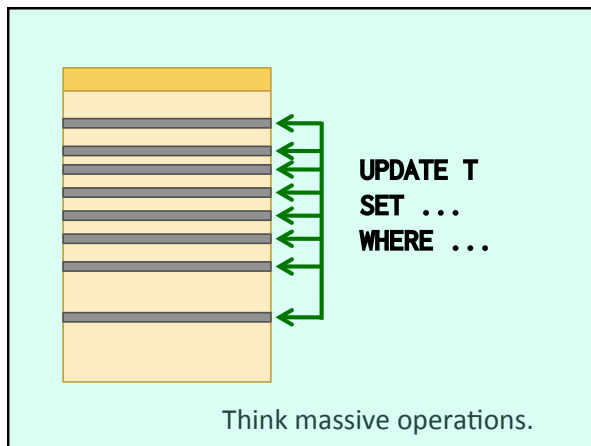
set

A very important point to remember is that UPDATE is a SET operation. The set may contain 1, or 1,000,000 or 0 rows.



Loop on SELECT  
UPDATE T  
SET ...  
WHERE KEY = ...


Updates in loops are WRONG (and very slow compared to the one-shot operation)



**Like a join in a select ...**

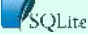
**... same issues with nulls and duplicates !**

update movies\_2  
set duration = (select duration  
from us\_movie\_info i  
where i.title = movies\_2.title  
and i.year\_released = movies\_2.year\_released),  
color = (select case color  
when 'C' then 'Y'  
when 'B' then 'N'  
end color  
from us\_movie\_info i  
where i.title = movies\_2.title  
and i.year\_released = movies\_2.year\_released)  
where country = 'us'  
and exists (select null  
from us\_movie\_info i2  
where i2.title = movies\_2.title  
and i2.year\_released = movies\_2.year\_released)

**Not found ?** 

**NULL**

As subqueries can return NULL, you must be certain to only affect rows in your scope.

 SQLite


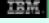
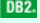
```

update movies_2
set duration = (select duration
                from us_movie_info i
                where i.title = movies_2.title
                  and i.year_released = movies_2.year_released),
  color = (select case color
                when 'C' then 'Y'
                when 'B' then 'N'
              end color
            from us_movie_info i
            where i.title = movies_2.title
              and i.year_released = movies_2.year_released)
where country = 'us'
  and exists (select null
              from us_movie_info i2
              where i2.title = movies_2.title
                and i2.year_released = movies_2.year_released)

```

**Three Queries per row processed**

Not madly efficient; all subqueries are correlated (for the third query SQLite now supports the same as Oracle).

 ORACLE  IBM  DB2

```

update movies_2
set (duration, color) =
  (select duration,
    case color
      when 'C' then 'Y'
      when 'B' then 'N'
    end color
   from us_movie_info i
   where i.title = movies_2.title
     and i.year_released = movies_2.year_released)
where country = 'us'
  and (m.title, m.year_released)
  in (select title, year_released
      from us_movie_info)

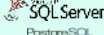

```

**Not found ?** ← run for each retrieved row

**NULL**

**Once**

Oracle and DB2 both support subqueries returning several columns (SQLite also now).

 SQL Server  PostgreSQL


```

update movies_2
set duration = i.duration,
  color = case i.color
    when 'C' then 'Y'
    when 'B' then 'N'
  end
from us_movie_info i
where i.title = movies_2.title
  and i.year_released = movies_2.year_released
  and movies_2.country = 'us'

```

**3 Join**

SQL Server and PostgreSQL both support the same older-join type of syntax allowing to join the updated table to the one from which we are getting data.

 MySQL

```

update movies_2 m
  inner join us_movie_info i
    on i.title = m.title
   and i.year_released = m.year_released
set m.duration = i.duration,
  m.color = case i.color
    when 'C' then 'Y'
    when 'B' then 'N'
  end
where m.country = 'us'

```

MySQL allows a join with the newer syntax.

What can happen when join conditions are

# WRONG?

When you have a SELECT wrong, it only affects your query. When you have an UPDATE wrong, you can corrupt the database and later correct queries on wrong data will return wrong results. So you really need to be extra careful.

**movies\_2**      **us\_movie\_info**

**title** ↔ **title**

**year\_released** ↔ **year\_released**

**country = 'us'**

Imagine for instance that we forget the join on the year and that we have remakes. What will happen?

**movies\_2**



**us\_movie\_info**



Let's first say that we only have remakes in the table that we update.

Running a SELECT shows what happens.

```
select m.title, m.year_released,
       i.year_released, i.duration, i.color
from movies_2 m
     inner join us_movie_info i
       on i.title = m.title
where m.title like 'Treasure%'
```

One row from the source table will be associated with both films.

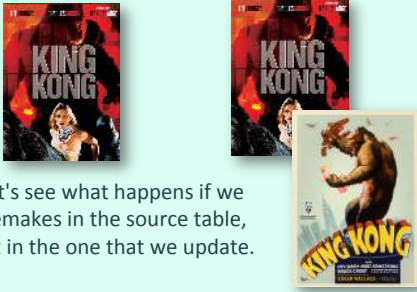
title	year_released	year_released	duration	color
Treasure Island	1934	1934	103	B
Treasure Island	1950	1934	103	B

**movies\_2**

**us\_movie\_info**

We'll set wrong information for the 1950 film.

**movies\_2**      **us\_movie\_info**



Now let's see what happens if we have remakes in the source table, but not in the one that we update.

Once again, a SELECT shows what happens.

```
select m.title, m.year_released,
       i.year_released, i.duration, i.color
from movies_2 m
     inner join us_movie_info i
       on i.title = m.title
    where m.title = 'King Kong'
```

The same row will be updated twice. What will remain is the last update. Heads or tails?

title	year_released	year_released	duration	color
King Kong	1976	1933	100	B
King Kong	1976	1976	134	C

**movies\_2**      **us\_movie\_info**

**Subquery**

```
update movies_2
set duration =
  (select duration
   from us_movie_info i
   where i.title = movies_2.title)
```

... Note that a subquery returning more than one row would generate an error.

**Join**

title	year_released	year_released	duration	color
King Kong	1976	1933	100	B
King Kong	1976	1976	134	C

A join won't fail, and just update randomly.

**2 rows** (with arrow pointing to the two rows in the table)


**FAILURE** (in orange)

**Do you feel**

**lucky?**

## Same Rules Apply for **UPDATE** as for **SELECT**

Except that as already stated, an update can change the data wrongly.

Surname	Birthdate	Picture
Hepburn	4-May-1929	

**STOP**  
Katharine

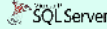

You are reminded that if a regular attribute can be updated, it's usually forbidden to update a key - it's the identifier. You cannot change an identifier. You can only delete the row and insert another.


### Primary Key



Update-wise, a primary key is locked.  
Off-limits.

Picture by Andrew Magill




### Update or Insert

A interesting operation would be to update a film we know, and insert it if we don't. That's the purpose of MERGE.

```
merge into movies_2 m
using (select 'us' as country,
            title,
            year_released,
            duration,
            case color
              when 'C' then 'Y'
              when 'B' then 'N'
            end as color
      from us_movie_info) i
on (i.country = m.country
and i.title = m.title
and i.year_released = m.year_released)
when matched then
  update
  set m.duration = i.duration,
      m.color = i.color
when not matched then
  insert(title, year_released, country, duration, color)
  values(i.title, i.year_released, i.country, i.duration, i.color)
```

*Oracle: movieid + sequence*




### Update or Insert

```

insert into movies_2(title, year_released,
                    country, duration, color)
select title, year_released, country, duration, color
from (select title,
            year_released,
            'us' as country,
            duration,
            case color
              when 'C' then 'Y'
              when 'B' then 'N'
            end color
      from us_movie_info) i
on duplicate key update
  movies_2.duration = i.duration,
  movies_2.color = i.color

```

MySQL can catch an insert that fails because the row is already here, and turn on the fly the insert into an update.



### Update or Insert

```

insert or replace into movies_2(title, year_released,
                               country, duration, color)
select title, year_released, country, duration, color
from (select title,
            year_released,
            'us' as country,
            duration,
            case color
              when 'C' then 'Y'
              when 'B' then 'N'
            end color
      from us_movie_info) i

```

SQLite allows something similar with a simpler (but less flexible) syntax. Beware, because it deletes a row and creates a new one, foreign keys may not like it.

~~Update or Insert~~  
<sup>then</sup> Update or Insert

**Update** +

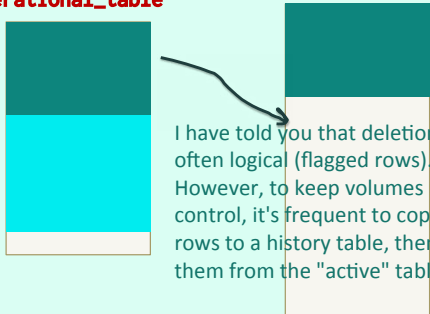
```

insert into movies_2(title, year_released, country,
                    duration, color)
select i.title, i.year_released, 'us', i.duration,
      case i.color
        when 'C' then 'Y'
        when 'B' then 'N'
      end
from us_movie_info i
left outer join movies_2 m
  on m.title = i.title
 and m.year_released = i.year_released
 and m.country = 'us'
where m.movieid is null

```

When none of the above is available, you should try to update, and if nothing is affected insert. NEVER count first to see if the row is already here! It's useless work.

**operational\_table**      **history\_table**



I have told you that deletions are often logical (flagged rows). However, to keep volumes under control, it's frequent to copy old rows to a history table, then delete them from the "active" table.



**delete from table\_name  
where ...**

If you omit the WHERE clause, then (as with UPDATE) the statement affects all rows and you

**Empty table\_name !**

But of course you NEVER work in autocommit mode and always execute a big update or delete in a transaction, don't you?

**rollback**

That's when you feel grateful for some features.

Flickr: Giò-S.p.o.t.s.



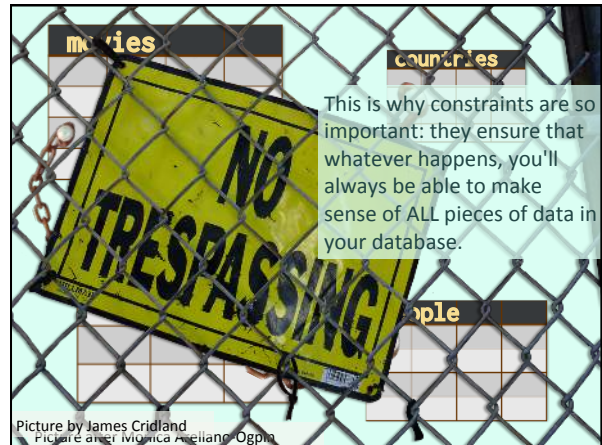
As DELETE saves data for rollback before removing it, it can be slow. There is a TRUNCATE (without a WHERE clause) that cannot be rolled back and is far more efficient. Leave it to senior DBAs with little remaining gray hair.

Constraints  
**= guarantee**

One important point with constraints (foreign keys in particular) is that they guarantee that data remains consistent. They don't only work with INSERT, but with UPDATE and DELETE as well.

## Try to delete rows from table **countries**

For instance, you can delete a country for which there are no movies. As soon as you have one movie, you are prevented from deleting the country otherwise the foreign key on table MOVIES would no longer work for films from that country.



Removing automatically dependencies kind of defeats the purpose of constraints. It's acceptable to delete "in your back" what was created "in your back" ( we are going to see triggers very soon)

A constraint can be created with  
on delete cascade  
or  
on delete set null  
**Rare!**

## begin transaction

Keep in mind that rows you change are locked for all the duration of a transaction: a clear consequence is that, usually, your transactions shouldn't last too long if several users are modifying data at the same time.



## commit

Flickr: Andrew Magill

## SQL: declarative

### What about programming?

SQL is essentially a declarative language: you state what you want, and the DBMS is supposed to manage by itself how to find it (the sad reality is that how you express yourself often makes a difference). How can we code with SQL?

Let's first see how you can code INSIDE the database.

## BUSINESS LOGIC DATA

In a business application, you have some business logic that must be applied to the data. We'll see later various options about when and where business logic can be applied, but you can do an awful lot inside the database.

Most DBMS (the exception is SQLite, not a true DBMS) implement a built-in, SQL-based programming language, that can be used when a declarative language is no longer enough. Let's start with the simplest thing, defining functions.

## Functions

In a business application, you have some business logic that must be applied to the data. We'll see later various options about when and where business logic can be applied, but you can do an awful lot inside the database.

first_name	surname
Erich	Stroheim (von)

I gave an update example in which I was modifying every name starting with 'von ' so that they sort properly.

```
select first_name || ' ' || surname as full_name
from people;
```


**Erich Stroheim (von)**

Sorting is one thing, but if I ever want to display the full name of a person by concatenating first\_name and surname, it will look weird for von Stroheim. What I really want to see is

Erich von Stroheim

```
case
  when first_name is null then ''
  else first_name || ' '
end
|| case position('(' in surname)
  when 0 then surname
  else trim(')' from substr(surname,
                           position('(' in surname) + 1))
  || ' '
  || trim(substr(surname, 1,
                 position('(' in surname) - 1))
end full_name
```

first_name	surname
Erich	Stroheim (von)
Erich von Stroheim	



Needless to say, whenever you have painfully written something as complicated, which is pretty generic, you'd rather not copy and paste the code every time you need it.

```
case
  when first_name is null then ''
  else first_name || ' '
end
|| case position('(' in surname)
  when 0 then surname
  else trim(')' from substr(surname,
                           position('(' in surname) + 1))
  || ' '
  || trim(substr(surname, 1,
                 position('(' in surname) - 1))
end full_name
```

Flickr:Kevin Rawlings

You'd like to store the expression and reuse it in another context. In fact you can.

# STORE FOR REUSE

Here is a PostgreSQL example.




```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end |
    case position('(' in p_sname)
        when 0 then p_sname
        else trim('(' from substr(p_sname,
                                position('(' in p_sname) + 1))
            || ' '
            || trim(substr(p_sname, 1,
                          position('(' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

```
select full_name(first_name, surname) as name,
       born, died
from people
order by surname
```

Once your function is created, you can use it as if it were any built-in function.

Note that you usually have to write your functions in the provided language for safety: a badly coded C function could take down a whole server, corrupt data, etc. The provided language provides a kind of sand-boxed environment.

## Procedural extensions to SQL



**T-SQL**



**(no name)**



**PL/SQL**



**PL/PGSQL**



**SQL PL**



**nothing ...**

You can use C or any language with SQLite. If you crash your program, it only affects you.

## Procedural ?

**variables**

**conditions**

**loops**

**arrays**

**error management**

## ... TRUE PROGRAMMING LANGUAGE

They also support all DML statements (no DDL, but you can cheat)

Procedural extensions provide all the bells and whistles of true programming languages (they were often inspired by programming languages such as PL/I or ADA). They are a mixed blessing, because they often incite programmers to do the wrong things with them.

```
select col1, col2, ...
into local_var1, local_var2, ...
from ...
```

## + CURSORS

To retrieve data from the database into your variables, you can use SELECT ... INTO ... if your query returns a single row, or you can use cursors, which are basically "row variables" that are used for iterating over what a query returns.



**Cultural mismatch**

**row-by-row**

**set processing**

And here we have a problem, because there is a big cultural gap between the relational mindset and procedural processing.

Flickr: Jeff Sandquist

## BAD EXAMPLE

In the category "never, ever do that even if you encounter it often" there is the infamous "look-up" function that returns for instance the label associated with a value. Because it's a procedure stored inside the database, many developers believe in good faith that's how things should be done. Definitely no.

```
ORACLE

create function country_name(p_code varchar2)
return countries.country_name%type
as
    v_name    countries.country_name%type;
begin
    select country_name
    into v_name
    from countries
    where country_code = p_code;
    return v_name;
end;
```

# NO