

CS209

Computer system design and application

Stéphane Faroult
faroult@sustc.edu.cn

Zhao Yao zhaoy6@sustc.edu.cn
Liu Zijian liuzijian47@163.com
Li Guansong intofor@163.com

URL

"http://edu.konagora.com/video/TestVid.mp4"
"file:///Users/roughsea/Movies/TestVid.mp4"

Media

We have seen last time that to display a video in JavaFx you need to provide a URL to the Media constructor, which refers to a file that may be local or distant. The Media constructor manages it all by itself, the first characters in the URL tells it how it can obtain the file – by requesting it from a distant HTTP server in the first case or from the local file system in the second one. Transparent.

MediaView

MediaPlayer

The Media is added to a MediaPlayer that provides methods for controlling the media (playing it, pausing it and so forth). The MediaPlayer is in turn added to a MediaView which is what JavaFx can display.

Media

I have in my demo application added to a "border Pane" the MediaView and a Hbox (Horizontal box) that contains widgets that call the MediaPlayer methods.

Hbox for controls

MediaView


I could have added more controls (see the MediaPlayer docs). The "Play" button morphs into a "Pause" button when clicked ("Pause" becomes in turn "Play" when clicked)

Skinning a GUI using CSS

Last JavaFx subject, superficial in all meanings of the word but important (looks sell): how to change the appearance of a JavaFx application. I have already briefly talked about it, the best way is to do it through an external style sheet (.css file). People will be able to change, often in a very impressive way, the looks of your application by changing this file and without any need to access the code (in fact, they just need the .css and the .class to run the "modified" application).

If you want to see how far you go with "styling", you can visit <http://csszengarden.com> and click on designs on the right hand-side. The same page will look completely different.

I have already said it, CSS stands for Cascading Style Sheet and Cascade is French for waterfall. The name emphasizes that styles "drip down", and are inherited from parent structures.



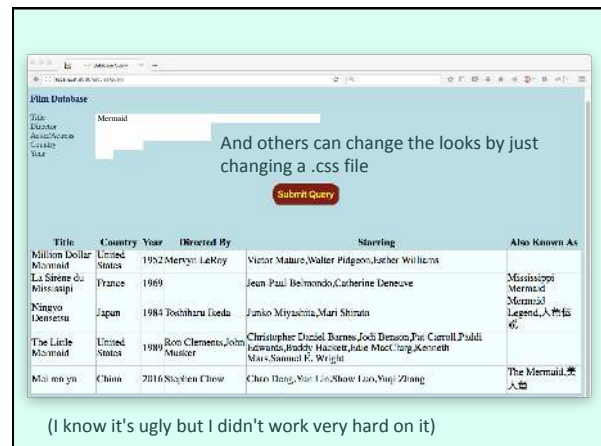
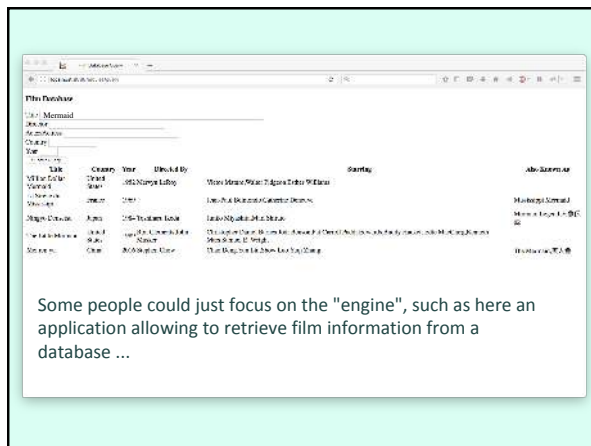
It all starts with HTML HyperTextMarkupLanguage

The idea is stolen from web applications. Web applications just send HTML pages to browsers that decode and display these pages. An HTML page is organized by sections between pairs of tags (<tag> at the beginning, </tag> at the end) that structure the document and can contain in turn other pairs of tags, thus defining a kind of hierarchical structure (note that some tags, such as those for images, act both as opening and closing tags). Tags pretty often also contain attributes (such as the image file name for an image tag).

<tags>

Associate formatting
with each tag in a
"stylesheet"

In the very early days of the web, people were using tags to format their pages, for instance what they wanted in bold was between **<bold>** and **</bold>** (inspired by previous document generation systems that were sending special signals to printers), and you could change the fonts with **<font** attributes specifying the font, size and everything> ... ****. As websites were growing in size and number of pages, and as increasingly pages were generated by programs instead of being created by hand, it became unmanageable, especially when the marketing department was deciding on new corporate colors. So the idea was to associate formatting to tags in one or several separate text files.

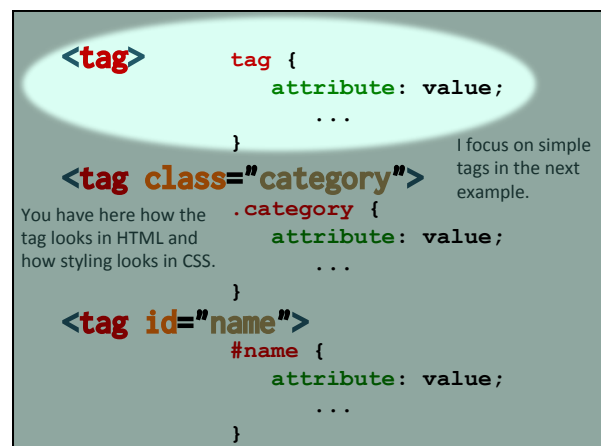


The way it works in web pages

Before I discuss about CSS in JavaFx, I'm going to talk about CSS with HTML, because there is far more CSS written for HTML than for JavaFX.

There are three main ways to specify how to display what is between tags.

1. You can specify for a given tag, associating a tagname with visual characteristics
2. I have mention that tags can have attributes, one is "class" (unrelated to object-oriented programming) listing one or several categories. This allows to create a subcategory, or to give some common visual characteristics across different tags that have the same class (for instance "inactive")
3. You can give another attribute "id" to a tag, an this allows you to make one particular tag look really special.



<body>

In a HTML page, everything that is displayed is between <body> and </body> tags, so <body> really defines the global looks of the page.

In my example, I'll use tags <a> associated with a link, <h1> with a (first level) header, <table>, <th> (table header) and <td> (table data). I have omitted <tr> (table row) which usually surrounds the columns of a same row. The tag will be there on the page, but I'm not associating any special style with it in my example.

Link <a>

Title <h1>

| Header1 | Header2 |
|---------|---------|
| blah | blah |
| blah | blah |

<table> <th> <td>

```
body {text-align: center;
      background-color: black;}
a {color: lightskyblue;}
a:visited {color: lightseagreen;}
a:hover {color: salmon;}
h1 {color: cornsilk;}
th {background-color: sienna;
     color: orange;}
td {background-color: burlywood;}
table {width: 80%;
       margin-left: auto;
       margin-right: auto;
       text-align: center;}
form {width: 50%;
      margin-left: auto;
      margin-right: auto;
      padding: 20px;
      background-color: silver;}
```

Link

Title

| Header1 | Header2 |
|---------|---------|
| blah | blah |
| blah | blah |

css color names

This will give you the names of colors understood by browsers.

Notice the special a:visited (link you have already clicked on) and a:hover (when the cursor moves over the link)

```
body {text-align: center;
      background-color: black;}
a {color: lightskyblue;}
a:visited {color: lightseagreen;}
a:hover {color: salmon;}
h1 {color: cornsilk;}
th {background-color: sienna;
     color: orange;}
td {background-color: burlywood;}
table {width: 80%;
       margin-left: auto;
       margin-right: auto;
       text-align: center;}
form {width: 50%;
      margin-left: auto;
      margin-right: auto;
      padding: 20px;
      background-color: silver;}
```

styles.css

Styling is written to a file that is included in the HTML page by a special tag in the <header>...</header> section that precedes the "body". Then you can change it when needed.

<link rel="stylesheet" href="styles.css"/>

The way it works in javafx

Nodes ^{almost} are equivalent to tags

.root plays the same role as **body**
otherwise use class names prefixed with a dot

.button

You have of course no tags in a JavaFx application, but you have the same kind of hierarchy through nodes, containers and widgets. The JavaFx class names are used with the same syntax as the HTML classes in CSS, prefixed by a dot.

The way it works in javaFX

Nodes are equivalent to tags

Attribute names are prefixed with **-fx-**

-fx-font-size: 150%;

CSS attributes also have a special name with JavaFx. The idea is to be able to have a single CSS files shared by a Web and a JavaFx application without having any conflict.

The way it works in javaFX

Nodes are equivalent to tags

Attribute names are prefixed with **-fx-**

`Node.setId("name")`

`Node.getStyleClass().add("css class")`

Finally, node methods allow you to associate with a node the same kind of attributes as with a HTML tag. You can only have a single Id, but you can have several classes, and therefore `getStyleClass()` returns a list.

The way it works in javaFX

Nodes are equivalent to tags

Attribute names are prefixed with **-fx-**

`Node.setId("name")`

`Node.getStyleClass().add("css class")`

`Node.setStyle("-fx-attribute: value")`

`Scene scene = new Scene(new Group(), 500, 400);`
`scene.getStylesheets().add("path/styles.css");`

We have already seen how to load the CSS file into the JavaFX application.

Everything cannot be styled with CSS in JavaFX

CSS styling allows you to go rather far in JavaFx, but not as far as you could go in HTML. Some elements may prove hard to style with CSS. You may sometimes have to code some styling in the Java application. However, if you still want this styling to be "externalized", don't forget that properties files also provide a way to read attributes at run-time. It's of course better to have all styling at one place, but it's better to use a properties file than to hard-code.



A Properties file might sometimes be a workaround

PERSISTENCE

After Graphical User Interfaces, our main "big" topic will be the topic of persistence.

Another word that comes from Latin, and conveys an idea of continuity and robustness.


Data Persistence



von Neumann
Architecture



In the von Neuman
architecture we have
everything in memory.





But what happens when the
power is cut? Poof, gone. We can
reload programs when we restart,
but if they modify data we'd
rather save the work done
somewhere.



save, keep
Need to **persist**
program results
somewhere



That's all the idea of persistence. "Somewhere" may take multiple shapes (including a remote computer)

Especially in the old days

Memory

Memory is like cars: fast goes hand-in-hand with expensive. Fast memory is "volatile", which means (that's the meaning of the word) that it's content can "fly away" (if you cut power). Slow memory stays, but the problem is that when you access it your program works at its slow speed.

volatile

non-volatile

Flickr: Andrew Magill

So we have to work as much as we can in memory, and only in memory, for speed ...

Mostly work in memory for speed


Memory

... while we still need a safety net.

Write to file for safety

magnetic


Non-Volatile Memory Technology



There are several non-volatile technologies available. The oldest one, still alive and kicking, is the magnetic disk, which has done tremendous progress over the years in data density and transfer rates (not so much directly accessing data somewhere on the disk, called "random access")

optical


Non-Volatile Memory Technology



Optical technologies are no longer so hot as they used to be, but they are very good for archiving and, because readers/writers are really mass produced, fairly cheap. Data updates? Better to look elsewhere.

solid-state

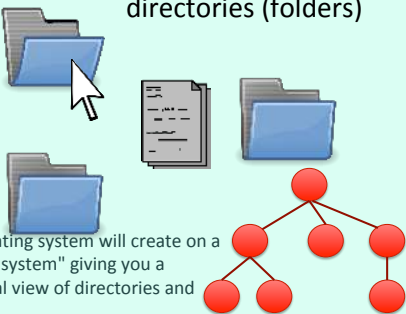
Non-Volatile Memory Technology



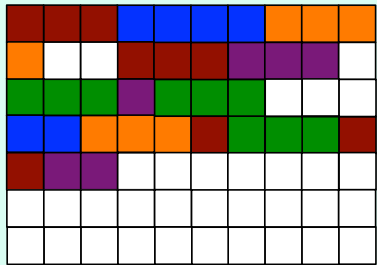
Solid-state technologies are very good for reading, much less so for writing.

File System

directories (folders)



Your operating system will create on a disk a "file system" giving you a hierarchical view of directories and files.



But in reality the system will reserve blocks, all multiples of a same unit, that it will associate with one file. Files will grow, deletions will create gaps soon filled by blocks from from other files, and the system will try to do all this as efficiently as possible. There is a strong disconnect between the view we have and physical reality.

Stream redirection

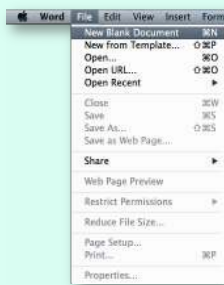
Instead of keyboard

```
$ java MyProgram < input_file
```

```
$ java MyProgram > output_file
```

Instead of screen

One very easy way to save data and to restore it is, instead of doing it in Java, to let the system do it through what is known as "Stream redirection". Your program may read from what it thinks is the keyboard when in reality input comes from a file, and what is written to the screen can also be redirected to a file.



Files may be handled directly by programs.

Many programs, though, will directly handle file either when you specifically request it, or in your back (for instance "autosave" every ten minutes).

Two traditional types of files:



Text files *← only printable characters*



Binary files

If there are tons of file extensions and many types of files, they all fall in one of two categories: text or binary. The only problem is, whatever they are, they are all made of 0s and 1s. So really it's all a matter of interpretation.

Interpretation is the big, hard question. You cannot guess the meaning of 1s and 0s just by looking at them. You must have an idea already. And even with text, there are many different ways to encode one single character (and don't believe that the problem doesn't exist even with basic Latin letters – there is another encoding system than ASCII called EBCDIC and the bits meaning 'a' in ASCII mean '/' in EBCDIC). If you haven't the key allowing you to decrypt the bits, you are lost.

HOW to understand the 0's and 1's?

Two traditional types of files:



Text files ← only printable characters
WHEN DECODING AS
CHARACTERS



Binary files

So the true definition of "text file" is that it only contains characters that you can print (including spaces and carriage returns) when you decrypt the file as a bunch of characters.

Text files

There are many types of text files – not only files with extension .txt !

Can be opened by a "text editor" (eg Notepad) or displayed using **more** (Linux) or **type** (Windows)

Program code (.c, .h, .py, .php, .java, .bat, .sh, ...)

Plain text (.txt, .ini)

Text with readable tags (.html, .rtf, .xml)

Data as text (.csv)

Only contain printable characters

Binary files

There are also many types of binary files, including those used by documents that are supposed to be mostly text ...

Can only be opened by a special program

Compiled program (.o, .exe, .class)

Archives, compressed files (.tar, .tgz, .zip)

Crypted files

Text with non readable formatting (.docx, .pdf, .xlsx, .pptx)

Multimedia (.gif, .jpg, .png, .mp3, .wav, .mpg, .flv, ...)

Binary files

Very often (but not always) binary files are a basic "dump" of what you have in memory. When the file may be written on one system and read on a very different one, some standard encoding may be applied.

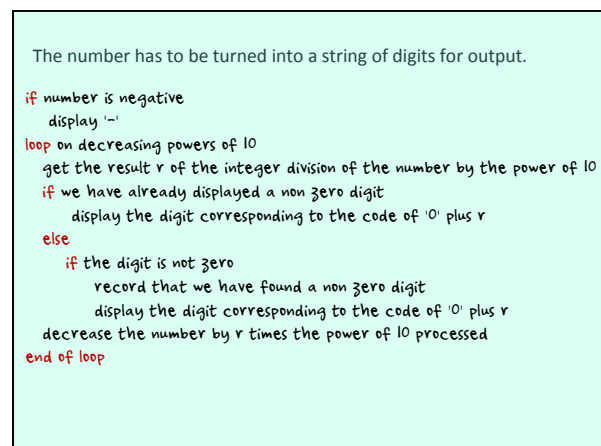
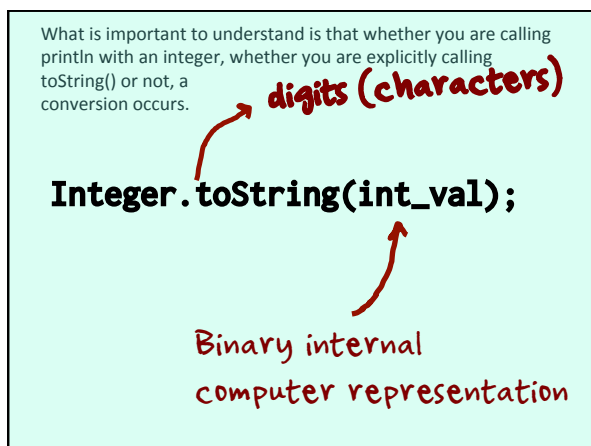
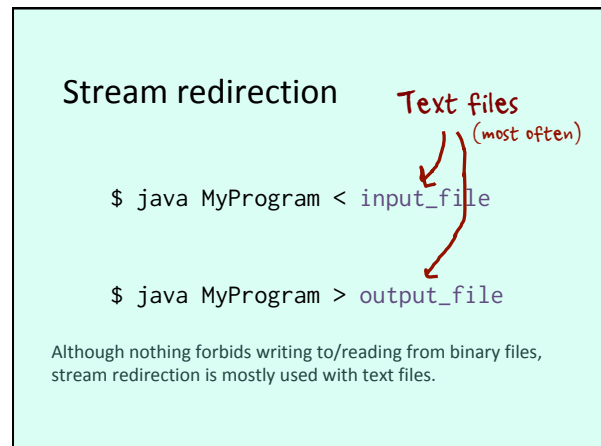
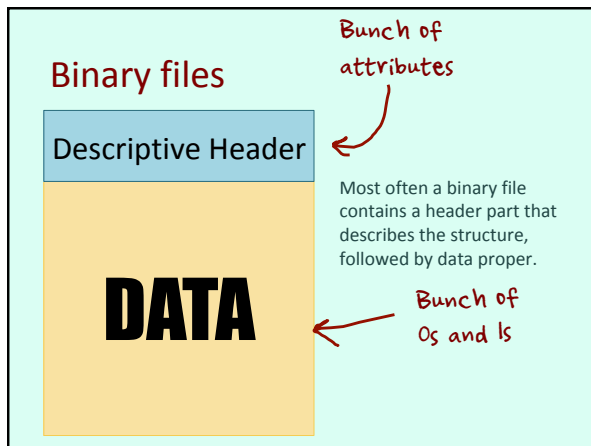
Memory structures written "as is"

More compact (no difference for text)

No conversion during I/Os

May be portability issues between computers (Windows/Mac/Linux)

One big problem is the "small endian"/"big endian" issue, which is a hardware issue. The 4 bits that make up half a byte may be swapped.



Input requires the opposite.

Integer.parseInt() or the method **nextInt()** of a **Scanner** object perform the reverse operation

HOW to understand the 0's and 1's?

```
public class Hello {
    public static void main(String[] args) {
        System.out.println("Hello!");
    }
}
```

Let's check **Hello.class**

A program has, to "understand a file", a number of options.

HOW to understand the 0's and 1's?

1

Assume that it's what we expect

for instance, text ...

The simple one is "when the only tool you have is a hammer, everything looks like a nail".

```
$ cat Hello.class
????4

<init>()VCodeLineNumberTablemain([Ljava/lang/String;)V
SourceFile
Hello.java

    Hello!
        ellojava/lang/Objectjava/lang/
SystemoutLjava/io/PrintStream;java/io/
PrintStreamprintln(Ljava/lang/String;)V!    *??

    %   ???

$ "cat" writes everything as text to the screen. The result looks
like garbage.
```

HOW to understand the 0's and 1's?

2 Assume that the extension is correct

I renamed Hello.class to Hello.c and tried to compile it. I got 105 warnings and 11 errors but the compiler tried.

HOW to understand the 0's and 1's?

3 Check the file header for a "magic number"

Binary files usually contain a "signature" in their header, a small number of bytes that are very specific to one type of files. You don't need to trust the extension.

```
$ od -x Hello.class
00000000 feca beba 0000 3400 1d00 000a 0006 090f
00000020 1000 1100 0008 0a12 1300 1400 0007 0715
00000040 1600 0001 3c06 6e69 7469 013e 0300 2928
00000060 0156 0400 6f43 6564 0001 4c0f 6e69 4e65
00000080 6d75 6562 5472 6261 656c 0001 6d04 6961
000000a0 016e 1600 5b28 6a4c 7661 2f61 616c 676e
000000c0 532f 7274 6e69 3b67 5629 0001 530a 756f
000000e0 6372 4665 6c69 0165 0a00 6548 6c6c 2e6f
00000100 616a 6176 000c 0007 0708 1700 000c 0018
00000120 0119 0600 6548 6c6c 216f 0007 0c1a 1b00
...
```

feca beba
cafe babe

od is a Unix command that dumps a file.
All .class files start with the same bytes.

Most file-related classes are in the java.io package (there is also a java.nio). Two types of streams, Character or Byte, which can all be buffered or unbuffered.

Handling files in Java

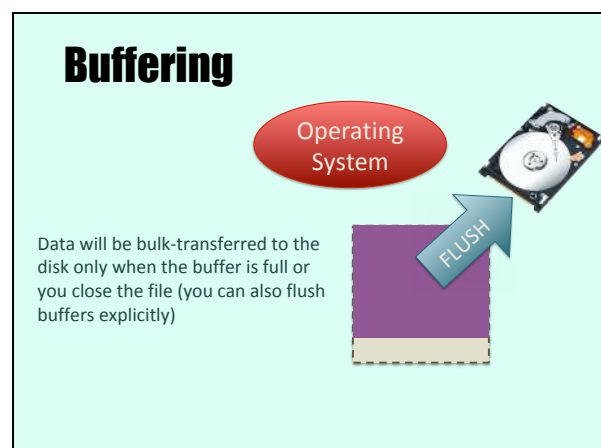
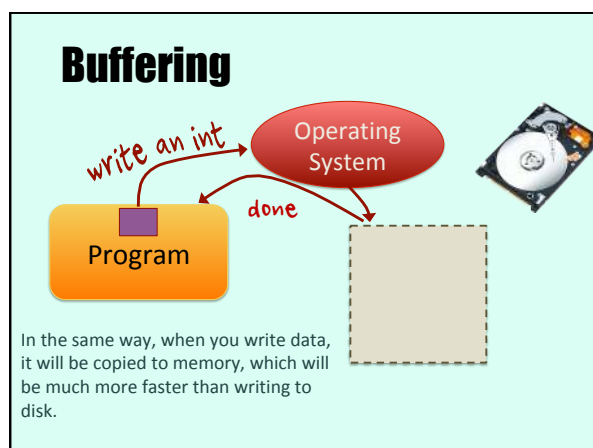
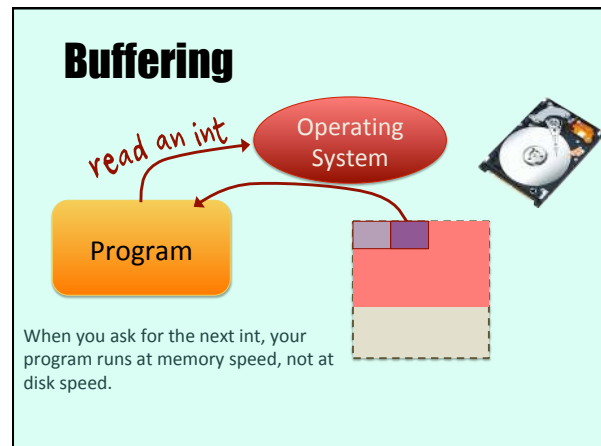
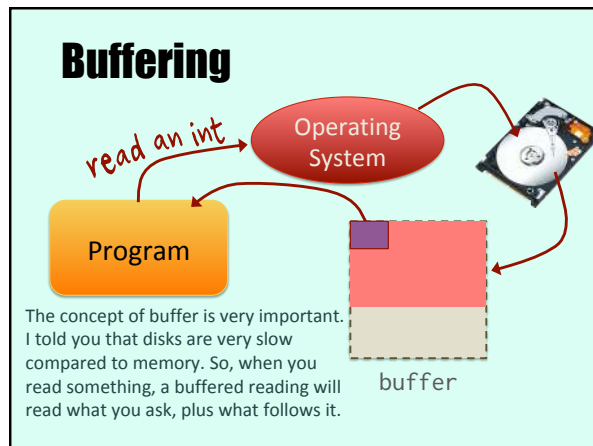
Key concept: **Stream**

Character streams

Buffered or **unbuffered**

Byte streams

```
import java.io.*
```




Buffering

What happens if the system restarts after a crash?

The answer of course is that what was in buffers is lost. Not a problem when reading, big problem when writing. It's not always easy to know what to replay.

No buffering



If you freak about crashes, you should use slow unbuffered operations.

Mostly used when writing critical information

Logs

Messaging

Performance - copying a 11M CSV file character by character

Test on my Mac (internal SSD)

| | |
|------------|-------------|
| Unbuffered | about 34.5s |
| Buffered | about 1s |

Test on my Mac (External USB HD)

| | |
|------------|------------|
| Unbuffered | about 36s |
| Buffered | about 1.4s |

How often does your computer crash? Would it be a complete disaster to run the program again after restart?


Performance - copying a 11M CSV file character by character

For 99% of cases, you should use buffered input/output operations (we could even say 100% for input).

Use **buffered** operations unless writing safely is a critical concern.

Or for debugging.

Sometimes hard to say what's going on today

Big disk systems have their own, battery protected, buffers (also called **CACHE**)  French for Hideout

Note that especially with high-end storage you rarely have one level of buffering (in which case unbuffered wouldn't be what it seems). It's a bit hard sometimes to know if the data is on disk or not, and the Cloud doesn't make it any simpler.

UNBUFFERED


```
InputStream in = null;
OutputStream out = null;
```

```
in = new InputStream(...);
out = new OutputStream(...);
```

The basis for all Input/Output operations are InputStreams and OutputStreams, which are unbuffered.

One thing that should not be forgotten with file operations is that it's probably the part of a program where everything can fail.

Lots of things can **GO WRONG**

Wrong file/directory name
Not allowed 
Content not as expected
Hardware problem (rare)

UNBUFFERED

```
InputStream in = null;
OutputStream out = null;
```

```
try {
    in = new InputStream(...);
    out = new OutputStream(...);
```

```
} catch ... {
} finally {
    if (in != null) {
        in.close();
    }
    if (out != null) {
        out.close();
    }
}
```

So you should really do everything in a try block, either a "try with resources" or a try with a "finally" block to make sure that files are cleanly closed and not left corrupt.

IMPORTANT!
Flushes everything and closes properly

UNBUFFERED

```

FileInputStream in = null;
FileOutputStream out = null;

try {
    in = new FileInputStream("filename");
    out = new FileOutputStream("filename");

} catch ... {
} finally {
    ...
}

```

File location is always a practical problem. Think of reflection and properties file.

Looks in the current directory unless you provide a full path

BUFFERED

```

BufferedInputStream in = null;
BufferedOutputStream out = null;

try {
    in = new BufferedInputStream(new InputStream(...));
    out = new BufferedOutputStream(new ...);

} catch ... {
} finally {
    ...
}

```

To turn an unbuffered stream into a buffered one, you just wrap the call to the stream constructor into a call to a buffered stream constructor.

InputStream and OutputStream are the parent classes for both byte and character streams.

Byte Streams

(not the most used)

There are libraries for multimedia

You may not use byte streams very often. Remember JavaFx: when you create a new Image or Media object, a binary file is read into memory by the constructor. There is necessarily a byte stream behind the scene, but it's all done by the constructor.

BYTE STREAM
(unbuffered)

```

FileInputStream in = null;
FileOutputStream out = null;

try {
    in = new FileInputStream("filename");
    out = new FileOutputStream("filename");

} catch ... {
} finally {
    ...
}

```

The examples I have previously shown are for byte streams ...

```

BufferedInputStream in = null;
BufferedOutputStream out = null;

try {
    in = new BufferedInputStream(new FileInputStream());
    out = new BufferedOutputStream(new ...);

} catch ... {
} finally {
    ...
}

```

... including the buffered version.

BYTE STREAM
(buffered)

CAUTION

When you talk about "bytes" in input/output operations, you are really dealing with **int** variables, not **byte** variables.

HISTORICAL REASON

Don't be misled by the "byte" in "byte stream". Operations deal with more than one byte.

Object Serialization

```

class Obj2 {
    private String name;
    private int value;
    ...
}

class Obj1 {
    private Obj2[] o = null;
    private int count = 0;

    public Obj1() {
        o = new Obj2[10];
    }
    ...
}

```

One application of byte streams is "serialization", dumping a memory object to file.

Object Serialization

When doing so you only want to store data, not memory addresses that will change when you reload.

```

class Obj2 {
    private String name;
    private int value;
    ...
}

class Obj1 {
    private Obj2[] o = null;
    private int count = 0;

    public Obj1() {
        o = new Obj2[10];
    }
    ...
}

```

You can also declare some attributes as "transient" to say they shouldn't be dumped.

| | | | |
|--------|---------|------|------|
| Amazon | 6992 | Nile | 6853 |
| 6853 | Yangtse | 6300 | |

Object Serialization: requisites

```

class Obj2 implements java.io.Serializable {
    private String name;
    private int value;
    ...
}

class Obj1 implements java.io.Serializable {
    private Obj2[] o = null;
    private int count = 0;
    public Obj1() {
        o = new Obj2[10];
    }
    ...
}

```

1 INTERFACE

To be able to do so you must first say that every object implements the Serializable interface.

NO method!

It's just a declaration, there is no method to implement (in fact, it just tells javac to generate what is needed)

Object Serialization: requisites

```

class Obj2 implements java.io.Serializable {
    private String name;
    private int value;
    ...
}

class Obj1 implements java.io.Serializable {
    private Obj2[] o = null;
    private int count = 0;
    public Obj1() {
        o = new Obj2[10];
    }
    ...
}

```

2 CONSTRUCTOR

YES!

A default constructor is helpful when recreating the object (although it looks that implementing Serializable seems to take care of that)

Object Serialization: requisites

And you need an "ObjectOutputStream" that is a special flavor of byte stream. This one comes by default with a buffer.

3 Stream

Buffer included

```

FileOutputStream fileOut = new FileOutputStream("file.dat");
ObjectOutputStream out = new ObjectOutputStream(fileOut);
out.writeObject(o);
out.close();
fileOut.close();

```

A file written on one computer can be read on any computer!

Now, I'm not impressed by performance. There may be cases when serialization performs very well, but it requires testing.

Character Streams

(most often)

Handle 16-bit unicode characters (**char** datatype)

You'll probably use Character Streams more often than byte streams. If you have a C background, don't forget that Java chars are 2 bytes, not one as in C.

CHARACTER STREAM
(unbuffered)

```

FileReader in = null;
FileWriter out = null;

try {
    in = new FileReader("filename");
    out = new FileWriter("filename");

    } catch ... {
    } finally {
        if (in != null) {
            in.close();
        }
        if (out != null) {
            out.close();
        }
    }
}

```

What used to be "Input" and "Output" with byte streams becomes "Reader" and "Writer" with character streams.

CHARACTER STREAM
(buffered)

```

BufferedReader in = null;
BufferedWriter out = null;

try {
    in = new BufferedReader(new FileReader(...));
    out = new BufferedWriter(new FileWriter(...));


    } catch ... {
    } finally {
        if (in != null) {
            in.close();
        }
        if (out != null) {
            out.close();
        }
    }
}

```

Otherwise it's exactly the same thing.

So what can character streams do that byte streams can't do?

Good question. In fact character streams have three features absent from byte streams.



1 They can read (write) lines

Need to use a **BufferedReader**

readLine() method

When buffered they can read or write a full line at once – because text files are usually a sequence of lines.

```

BufferedReader in = null;
String line;

try {
    in = new BufferedReader(new FileReader("..."));
    while ((line = in.readLine()) != null) {
        ...
    }
} finally {
    if (in != null) {
        try {
            in.close();
        } catch (IOException e) {
        }
    }
}

```

But beware of lines when text files were written on a system and are read on a different system.

%n in Java format

Which brings the interesting problem of lines. You probably know the carriage return character, '\n'. Java prefers '%n', which may be one or two characters depending on the system it runs on.

It all dates back to the glorious days of the typewriter, in which letters were always typed at the same place. It's the "carriage" around which the sheet was wrapped that moved from right to left.

Where is the Life we have lost in living? ↵
 Where is the wisdom we have lost in knowledge? ↵
 Where is the knowledge we have lost in information? ↵

TS Eliot ↵
 (1888-1965)

What was happening at the end of the line? You needed to scroll the paper (line feed) and push back the carriage to the right (carriage return).

Guess what, the first printers were computer-controlled typewriters (sort of).



Where everything becomes fun, it's that a Unix system separates lines with a single 'carriage return' character.

Where is the Life we have lost in living? ↵ Where is the wisdom we have lost in knowledge? ↵ Where is the knowledge we have lost in information? ↵ TS Eliot ↵
 (1888-1965)








\n

0x0A

00001010

While a Windows system still stores the two separate commands that used to be sent to the printer - line feed (represented by `\r`), then carriage return.

Where is the Life we have lost in living?  Where is the wisdom we have lost in knowledge?  Where is the knowledge we have lost in information?   TS Eliot 
(1888-1965)








`\r\n`

`0x0D0A`





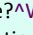



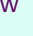

`0000110100001010`

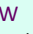

Linux file in a Windows editor

Where is the Life we have lost in living?  Where is the wisdom we have lost in knowledge?  Where is the knowledge we have lost in information?   TS Eliot 
(1888-1965)

A basic Windows editor (as Linux became more common, many editors became cleverer), looking for `\r\n` as separator between lines, will see in a Linux text file only one big line with `\n` characters that it doesn't know how to represent and that it will replace with squares.

Windows file in a Linux editor

Where is the Life we have lost in living?    
Where is the wisdom we have lost in knowledge?  
Where is the knowledge we have lost in information?    

TS Eliot 
(1888-1965) 

While on Linux the editor will understand correctly the `\n` in the Windows file, but will not know what to do with `\r` and will show it as `^W`. There is nothing simple in IT.

*There are conversion programs
Many programs understand both.*

MOST character files
are organized in lines
(variable or constant
length)

BUT a big character file
can sometimes contain a
single line (HTML, JSON,
XML ...)

You would of course
be wrong to believe
that a big text file
always contains
many lines.

Another interesting characteristic is that the parent class of `FileReader` allows you to specify the character encoding.

You can specify the character set in the constructor

2

They can change encoding

`InputStreamReader`

`FileReader`

And finally you can use a `Scanner` object with a character stream, which is very good at parsing text input.

3

You can use a **Scanner**

Very similar to keyboard and screen, which are character devices

File and Directory Operations

Check the **File** class

Copying, deleting files

Listing and searching directories

and so forth.

You can do a lot of things with files other than reading and writing them.

Files USED to be very important.

They still are, to some extents, for specialized applications, and they are still the backbone of persistence. However, as an application developers, you are increasingly isolated from files. I have mentioned the case of Image and Media; I could add Properties, and what we'll see the next times, databases that act as a layer between programs and files. A lot of data comes from networks as well.


Every application, 40 years ago, used to open and close a lot of files, this is no longer the case. You open files mostly to load data in memory, and work there. All this is related to the cost of memory.

1 Mb of memory

1970 \$730,000

When Dennis Ritchie created C, memory was horribly expensive. You had to release memory as soon as you no longer needed it, and try to save bytes.

C





Source : www.jcmit.com

1 Mb of memory

1990 \$100


Twenty years later, James Gosling could take a more relaxed approach, have a garbage collector freeing memory once in a while, and afford the luxury of a 2-byte char.



Source : www.jcmit.com

1 Mb of memory

2017



And today? Memory costs next to nothing. Just one problem: as the cost of memory was decreasing, applications were using more and more of it. And computers were supporting more and more users. You'd be wrong to believe that you no longer have to worry about memory. In some languages and environments (I'm thinking of Web servers running PHP) memory-per-user is limited to keep everything under control. But you don't need to fear loading sometimes quite a lot of data in memory.

Source : www.jcmit.com